

In [1]:

```

1 #ionosphere
2 import pandas as pd
3 import numpy as np
4 from sklearn.linear_model import LogisticRegression
5 from sklearn.preprocessing import StandardScaler

```

In [2]:

```

1 df=pd.read_csv(r"C:\Users\samit\Downloads\archive.zip")
2 df

```

Out[2]:

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.1	0.03760	...
0	1	0	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.00000	-0.04549	...
1	1	0	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.88965	0.01198	...
2	1	0	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.00000	0.00000	...
3	1	0	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.77152	-0.16399	...
4	1	0	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.14706	0.06637	...
...
345	1	0	0.83508	0.08298	0.73739	-0.14706	0.84349	-0.05567	0.90441	-0.04622	...
346	1	0	0.95113	0.00419	0.95183	-0.02723	0.93438	-0.01920	0.94590	0.01606	...
347	1	0	0.94701	-0.00034	0.93207	-0.03227	0.95177	-0.03431	0.95584	0.02446	...
348	1	0	0.90608	-0.01657	0.98122	-0.01989	0.95691	-0.03646	0.85746	0.00110	...
349	1	0	0.84710	0.13533	0.73638	-0.06151	0.87873	0.08260	0.88928	-0.09139	...

350 rows × 35 columns

In [3]:

```

1 pd.set_option('display.max_rows',10000000000)
2 pd.set_option('display.max_columns',10000000000)
3 pd.set_option('display.width',95)

```

In [4]:

```

1 print('This DataFrame has %d Rows and %d Columns'%(df.shape))
2 pd.set_option('display.width',95)

```

This DataFrame has 350 Rows and 35 Columns

In [5]:

```
1 df.head()
```

Out[5]:

	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.1	0.03760	0.85243
0	1	0	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.00000	-0.04549	0.5087
1	1	0	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.88965	0.01198	0.7308
2	1	0	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.00000	0.00000	0.0000
3	1	0	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.77152	-0.16399	0.5279
4	1	0	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.14706	0.06637	0.0378

In [6]:

```
1 features_matrix=df.iloc[:,0:34]
```

Type *Markdown* and LaTeX: α^2

In [7]:

```
1 target_vector=df.iloc[:,-1]
```

In [8]:

```
1 print('The features matrix Has %d Rows and %d Column(s)'%(features_matrix.shape))
```

The features matrix Has 350 Rows and 34 Column(s)

In [9]:

```
1 print('The target matrix Has %d Rows and %d Column(s)%(np.array(target_vector).res
```

The target matrix Has 350 Rows and 1 Column(s)

In [10]:

```
1 features_matrix_standardized=StandardScaler().fit_transform(features_matrix)
```

In [11]:

```
1 algorithm=LogisticRegression(penalty=None,dual=False,tol=1e-4,C=1.0,fit_intercept=T
```

In [12]:



```
1 Logistic_Regression_Model=algorithm.fit(features_matrix_standardized,target_vector)
```

In [13]:



```
1 observation=[[1,0,0.99539,-0.05889,0.8524299999999999,0.02306,
2              0.8339799999999999,-0.37708,1.0,0.0376,
3              0.8524299999999999,-0.17755,0.59755,-0.44945,0.60536,
4              -0.38223,0.8435600000000001,-0.38542,0.58212,-0.32192,
5              0.56971,-0.29674,0.36946,-0.47357,0.56811,-0.51171,
6              0.41070000000000003,-0.46160000000000003,0.21266,-0.3409,
7              0.42267,-0.54487,0.18641,-0.453]]
```

In [14]:



```
1 predictions=Logistic_Regression_Model.predict(observation)
2 print('The Model Predicted The Observation To Belong To Class %s'%(predictions))
```

The Model Predicted The Observation To Belong To Class ['g']

In [15]:



```
1 print('The Algorithm Was Trained To Predict One Of The Two Classes:%s'%(algorithm.c
```

The Algorithm Was Trained To Predict One Of The Two Classes:['b' 'g']

In [16]:



```
1 print("""The Model Says The Probability Of The Observation We Passed Belonging To C
2       (algorithm.predict_proba(observation)[0][0]))
3 print()
4 print("""The Model Says The Probability Of The Observation We Passed Belonging To C
5       (algorithm.predict_proba(observation)[0][1]))
```

The Model Says The Probability Of The Observation We Passed Belonging To Class['b'] is 3.317282690573631e-05

The Model Says The Probability Of The Observation We Passed Belonging To Class['g'] is 0.9999668271730943

In [17]:

```

1 #2
2 import re
3 from sklearn.datasets import load_digits
4 from sklearn.model_selection import train_test_split
5 import numpy as np
6 import matplotlib.pyplot as plt
7 import seaborn as sns
8 from sklearn import metrics
9 %matplotlib inline
10 digits=load_digits()

```

In [18]:

```

1 print("Image Data Shape",digits.data.shape)
2 print(" Label Data Shape",digits.target.shape)

```


Image Data Shape (1797, 64)
 Label Data Shape (1797,)

In [19]:

```

1 plt.figure(figsize=(20,4))
2 for index,(image,label)in enumerate(zip(digits.data[0:5],digits.target[0:5])):
3     plt.subplot(1,5,index+1)
4     plt.imshow(np.reshape(image,(8,8)),cmap=plt.cm.gray)
5     plt.title('Training:%i\n'%label,fontsize=10)

```



In [20]:

```

1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(digits.data,digits.target,test_size=

```

In [21]:

```
1 print(x_train.shape)
```

(1257, 64)

In [22]:

```
1 print(y_train.shape)
```

(1257,)

In [23]:



```
1 print(x_test.shape)
```

```
(540, 64)
```

In [24]:



```
1 print(y_test.shape)
```

```
(540,)
```

In [25]:



```
1 from sklearn.linear_model import LogisticRegression
2 logisticRegr=LogisticRegression(max_iter=1000)
3 logisticRegr.fit(x_train,y_train)
4 print(logisticRegr.predict(x_test))
```

```
[4 0 9 1 8 7 1 5 1 6 6 7 6 1 5 5 8 6 2 7 4 6 4 1 5 2 9 5 4 6 5 6 3 4 0 9
9
8 4 6 8 8 5 7 9 8 9 6 1 7 0 1 9 7 3 3 1 8 8 8 9 8 5 8 4 9 3 5 8 4 3 1 3
8
7 3 3 0 8 7 2 8 5 3 8 7 6 4 6 2 2 0 1 1 5 3 5 7 1 8 2 2 6 4 6 7 3 7 3 9
4
7 0 3 5 4 5 0 3 9 2 7 3 2 0 8 1 9 2 1 5 1 0 3 4 3 0 8 3 2 2 7 3 1 6 7 2
8
3 1 1 6 4 8 2 1 8 4 1 3 1 1 9 5 4 8 7 4 8 9 5 7 6 9 4 0 4 0 0 9 0 6 5 8
8
3 7 9 2 0 8 2 7 3 0 2 1 9 2 7 0 6 9 3 1 1 3 5 2 5 5 2 1 2 9 4 6 5 5 5 9
7
1 5 9 6 3 7 1 7 5 1 7 2 7 5 5 4 8 6 6 2 8 7 3 7 8 0 9 5 7 4 3 4 1 0 3 3
5
4 1 3 1 2 5 1 4 0 3 1 5 5 7 4 0 1 0 9 5 5 5 4 0 1 8 6 2 1 1 1 7 9 6 7 9
7
0 4 9 6 9 2 7 2 1 0 8 2 8 6 5 7 8 4 5 7 8 6 4 2 6 9 3 0 0 8 0 6 6 7 1 4
5
6 9 7 2 8 5 1 2 4 1 8 8 7 6 0 8 0 6 1 5 7 8 0 4 1 4 5 9 2 2 3 9 1 3 9 3
2
8 0 6 5 6 2 5 2 3 2 6 1 0 7 6 0 6 2 7 0 3 2 4 2 3 6 9 7 7 0 3 5 4 1 2 2
1
2 7 7 0 4 9 8 5 6 1 6 5 2 0 8 2 4 3 3 2 9 3 8 9 9 5 9 0 3 4 7 9 8 5 7 5
0
5 3 5 0 2 7 3 0 4 3 6 6 1 9 6 3 4 6 4 6 7 2 7 6 3 0 3 0 1 3 6 1 0 4 3 8
4
3 3 4 8 6 9 6 3 3 0 5 7 8 9 1 5 3 2 5 1 7 6 0 6 9 5 2 4 4 7 2 0 5 6 2 0
8
4 4 4 7 1 0 4 1 9 2 1 3 0 5 3 9 8 2 6 0 0 4]
```

C:\Users\samit\AppData\Local\Programs\Python\Python311\Lib\site-packages
 \sklearn\linear_model_logistic.py:458: ConvergenceWarning: lbfgs failed
 to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown i
 n:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

In [26]:

```
1 score=logisticRegr.score(x_test,y_test)
2 print(score)
```

0.9537037037037037

In [27]:

```
1 #gender submission
2 import pandas as pd
3 import numpy as np
4 from sklearn.linear_model import LogisticRegression
5 from sklearn.preprocessing import StandardScaler
```

In [28]:

```
1 df=pd.read_csv(r"C:\Users\samit\OneDrive\Desktop\jupyter\gender_submission.csv")
2 df
```

Out[28]:

	PassengerId	Survived
0	892	0
1	893	1
2	894	0
3	895	0
4	896	1
5	897	0
6	898	1
7	899	0
8	900	1
9	901	0

In [29]:

```
1 pd.set_option('display.max_rows',10000000000)
2 pd.set_option('display.max_columns',10000000000)
3 pd.set_option('display.width',95)
```

In [30]:

```
1 print('This DataFrame has %d Rows and %d Columns'%(df.shape))
2 pd.set_option('display.width',95)
```

This DataFrame has 418 Rows and 2 Columns

In [31]:

```
1 df.head()
```

Out[31]:

	PassengerId	Survived
0	892	0
1	893	1
2	894	0
3	895	0
4	896	1

In [32]:

```
1 features_matrix=df.iloc[:,0:34]
```

In [33]:

```
1 target_vector=df.iloc[:,-1]
```

In [34]:

```
1 print('The features matrix Has %d Rows and %d Column(s)'%(features_matrix.shape))
```

The features matrix Has 418 Rows and 2 Column(s)

In [35]:

```
1 print('The target matrix Has %d Rows and %d Column(s)%(np.array(target_vector).res
```

The target matrix Has 418 Rows and 1 Column(s)

In [36]:

```
1 features_matrix_standardized=StandardScaler().fit_transform(features_matrix)
```

In [37]:

```
1 algorithm=LogisticRegression(penalty=None,dual=False,tol=1e-4,C=1.0,fit_intercept=T
```

In [38]:

```
1 Logistic_Regression_Model=algorithm.fit(features_matrix_standardized,target_vector)
```


In [39]:

```
1 observation=[[1,0,]]
```

In [40]:

```
1 predictions=Logistic_Regression_Model.predict(observation)
2 print('The Model Predicted The Observation To Belong To Class %s'%(predictions))
```

The Model Predicted The Observation To Belong To Class [0]

In [41]:

```
1 print('The Algorithm Was Trained To Predict One Of The Two Classes:%s'%(algorithm.c
```

The Algorithm Was Trained To Predict One Of The Two Classes:[0 1]

In [42]:

```
1 print("""The Model Says The Probability Of The Observation We Passed Belonging To C
2       (algorithm.predict_proba(observation)[0][0]))
3 print()
4 print("""The Model Says The Probability Of The Observation We Passed Belonging To C
5       (algorithm.predict_proba(observation)[0][1]))
```

The Model Says The Probability Of The Observation We Passed Belonging To Class['b'] is 0.9855942924591363

The Model Says The Probability Of The Observation We Passed Belonging To Class['g'] is 0.014405707540863693

In [43]:

```
1 #train gender
2 import numpy as np
3 import pandas as pd
4 from sklearn import preprocessing
5 import matplotlib.pyplot as plt
6 plt.rc("font",size=14)
7 import seaborn as sns
8 sns.set(style="white")#while background style for seaborn plots
9 sns.set(style="whitegrid",color_codes=True)
10
11 import warnings
12 warnings.simplefilter(action='ignore')
```

In [44]:

```
1 train_df=pd.read_csv(r"C:\Users\samit\OneDrive\Desktop\jupyter\train.gender_submiss
2 train_df
3 test_df=pd.read_csv(r"C:\Users\samit\OneDrive\Desktop\jupyter\train.gender_submissi
4 test_df
```

Out[44]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.00	1	0	A/5 21171	7.2500
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.00	1	0	PC 17599	71.2833
2	3	1	3	Heikkinen, Miss. Laina	female	26.00	0	0	STON/O2. 3101282	7.9250
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.00	1	0	113803	53.1000
4	5	0	3	Allen, Mr. William Henry	male	35.00	0	0	373450	8.0500

In [45]:

```
1 train_df.head()
```

Out[45]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500

In [46]:



```
1 train_df.shape
```

Out[46]:

(891, 12)

In [47]:



```
1 test_df.head()
```

Out[47]:

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500



In [48]:



```
1 test_df.shape
```

Out[48]:

(891, 12)

In [49]:

```
1 train_df.describe()
```

Out[49]:

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

In [50]:

```
1 train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    int64
2   Pclass       891 non-null    int64
3   Name         891 non-null    object
4   Sex          891 non-null    object
5   Age          714 non-null    float64
6   SibSp        891 non-null    int64
7   Parch        891 non-null    int64
8   Ticket       891 non-null    object
9   Fare         891 non-null    float64
10  Cabin        204 non-null    object
11  Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

In [51]:



```
1 test_df.describe
```

Out[51]:

<bound method NDFrame.describe of

	PassengerId	Survived	Pclass
Name			
0	1	0	3
Mr. Owen Harris \			Braund,
1	2	1	1
Cumings, Mrs. John Bradley (Floren			
ce Briggs Th...			
2	3	1	3
en, Miss. Laina			Heikkin
3	4	1	1
Futrelle, Mrs. Jacques Heath			
(Lily May Peel)			
4	5	0	3
Allen, M			
r. William Henry			
5	6	0	3
M			
oran, Mr. James			
6	7	0	1
McCarthy, Mr. Timothy J			
7	8	0	3
Palsson, Maste			

In [52]:



```
1 test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    int64
2   Pclass       891 non-null    int64
3   Name         891 non-null    object
4   Sex          891 non-null    object
5   Age         714 non-null    float64
6   SibSp        891 non-null    int64
7   Parch        891 non-null    int64
8   Ticket       891 non-null    object
9   Fare         891 non-null    float64
10  Cabin        204 non-null    object
11  Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

In [53]:



```
1 train_df.isnull().sum()
```

Out[53]:

```
PassengerId    0
Survived        0
Pclass          0
Name            0
Sex             0
Age            177
SibSp           0
Parch           0
Ticket          0
Fare            0
Cabin          687
Embarked        2
dtype: int64
```

In [54]:



```
1 test_df.isnull().sum()
```

Out[54]:


```
PassengerId    0
Survived        0
Pclass          0
Name            0
Sex             0
Age            177
SibSp           0
Parch           0
Ticket          0
Fare            0
Cabin          687
Embarked        2
dtype: int64
```

In [55]:

```
1 ax=train_df["Age"].hist(bins=15,density=True,stacked=True,color='cyan',alpha=0.6)
2 train_df["Age"].plot(kind='density',color='teal')
3 ax.set(xlabel='Age')
4 plt.xlim(-10,85)
```

Out[55]:

(-10.0, 85.0)



In [56]:

```
1 print(train_df["Age"].mean(skipna=True))
2 print(train_df["Age"].median(skipna=True))
```

29.69911764705882

28.0

In [57]:

```
1 print((train_df['Cabin'].isnull().sum()/train_df.shape[0])*100)
```

77.10437710437711

In [58]:

```
1 print((train_df['Embarked'].isnull().sum()/train_df.shape[0])*100)
```

0.22446689113355783

In [59]:

```
1 print('Boarded passengers grouped by port of embarkation(C=Cherbourg,Q=Queenstown,S=Southampton):')
```

Boarded passengers grouped by port of embarkation(C=Cherbourg,Q=Queenstown,S=Southampton):


In [60]:

```
1 print(train_df['Embarked'].value_counts())
```

```
Embarked
S      644
C      168
Q       77
Name: count, dtype: int64
```

In [61]:

```
1 sns.countplot(x='Embarked',data=train_df,palette='Set2')
2 plt.show()
```



In [62]:



```
1 print(train_df['Embarked'].value_counts().idxmax())
```

S

In [63]:



```
1 train_data=train_df.copy()
2 train_data["Age"].fillna(train_df["Age"].median(skipna=True),inplace=True)
3 train_data["Embarked"].fillna(train_df["Embarked"].value_counts().idxmax(),inplace=True)
4 train_data.drop('Cabin',axis=1,inplace=True)
```

In [64]:



```
1 train_data.isnull().sum()
```

Out[64]:

```
PassengerId    0
Survived        0
Pclass         0
Name           0
Sex            0
Age            0
SibSp          0
Parch          0
Ticket         0
Fare           0
Embarked       0
dtype: int64
```

In [65]:



```
1 train_data.head()
```

Out[65]:

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare
0	1	0	3Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500
1	2	1	1Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833
2	3	1	3Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250
3	4	1	1Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000
4	5	0	3Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500

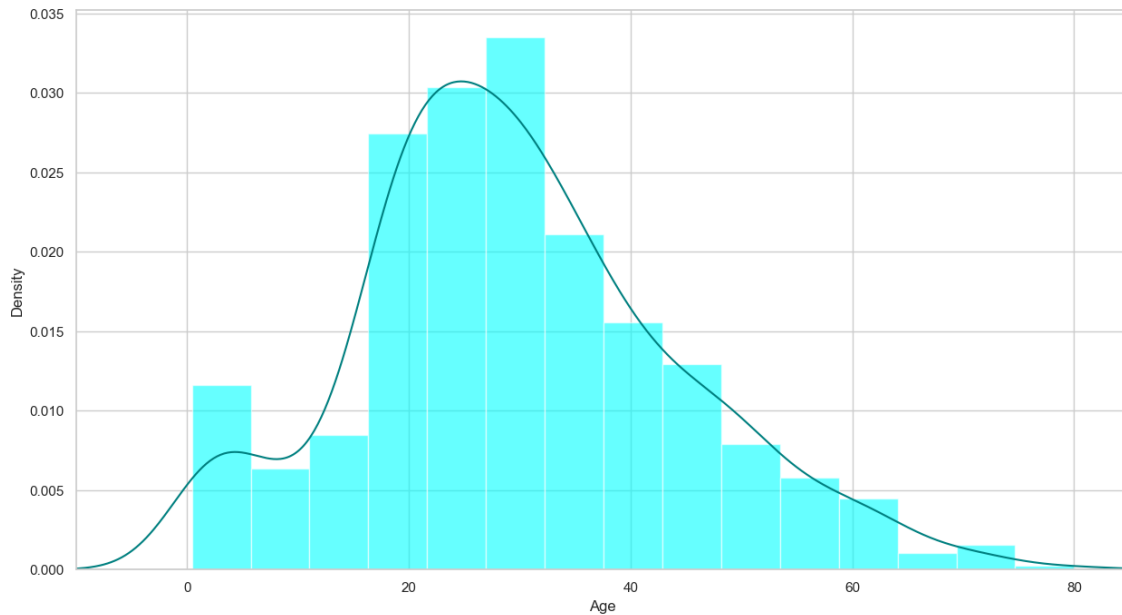


In [66]:

```
1 plt.figure(figsize=(15,8))
2 ax=train_df["Age"].hist(bins=15,density=True,stacked=True,color='cyan',alpha=0.6)
3 train_df["Age"].plot(kind='density',color='teal')
4 ax.set(xlabel='Age')
5 plt.xlim(-10,85)
6 plt.show
```

Out[66]:

<function matplotlib.pyplot.show(close=None, block=None)>



In [68]:

```
1 train_data['TravelAlone']=np.where((train_data["SibSp"]+train_data["Parch"])>0,0,1)
2 train_data.drop('SibSp',axis=1,inplace=True)
3 train_data.drop('Parch',axis=1,inplace=True)
```

In [72]:



```

1 training=pd.get_dummies(train_data,columns=["Pclass","Embarked","Sex"])
2 training.drop('Sex_female',axis=1,inplace=True)
3 training.drop('PassengerId',axis=1,inplace=True)
4 training.drop('Name',axis=1,inplace=True)
5 training.drop('Ticket',axis=1,inplace=True)
6
7 final_train=training
8 final_train.head()

```

Out[72]:

	Survived	Age	Fare	TravelAlone	Pclass_1	Pclass_2	Pclass_3	Embarked_C	Embarked
0	0	22.0	7.2500	0	False	False	True	False	
1	1	38.0	71.2833	0	True	False	False	True	
2	1	26.0	7.9250	1	False	False	True	False	
3	1	35.0	53.1000	0	True	False	False	False	
4	0	35.0	8.0500	1	False	False	True	False	



In [73]:



```
1 test_df.isnull().sum()
```

Out[73]:

```

PassengerId    0
Survived        0
Pclass          0
Name            0
Sex             0
Age            177
SibSp           0
Parch           0
Ticket         0
Fare            0
Cabin          687
Embarked        2
dtype: int64

```

In [76]:

▶▶

```
1 test_data=test_df.copy()
2 test_data["Age"].fillna(train_df["Age"].median(skipna=True),inplace=True)
3 test_data["Fare"].fillna(train_df["Fare"].median(skipna=True),inplace=True)
4 test_data.drop('Cabin',axis=1,inplace=True)
5
6 test_data['TravelAlone']=np.where((test_data["SibSp"]+test_data["Parch"])>0,0,1)
7
8 test_data.drop('SibSp',axis=1,inplace=True)
9 test_data.drop('Parch',axis=1,inplace=True)
10
11 testing=pd.get_dummies(train_data,columns=["Pclass","Embarked","Sex"])
12 testing.drop('Sex_female',axis=1,inplace=True)
13 testing.drop('PassengerId',axis=1,inplace=True)
14 testing.drop('Name',axis=1,inplace=True)
15 testing.drop('Ticket',axis=1,inplace=True)
```


In [77]:

▶▶

```
1 final_test=testing
2 final_test.head()
```

Out[77]:

	Survived	Age	Fare	TravelAlone	Pclass_1	Pclass_2	Pclass_3	Embarked_C	Embarked_S
0	0	22.0	7.2500	0	False	False	True	False	
1	1	38.0	71.2833	0	True	False	False	True	
2	1	26.0	7.9250	1	False	False	True	False	
3	1	35.0	53.1000	0	True	False	False	False	
4	0	35.0	8.0500	1	False	False	True	False	




In [80]:

```

1 plt.figure(figsize=(15,8))
2 ax = sns.kdeplot(final_train["Age"][final_train.Survived == 1], color="darkturquoise", shade=True)
3 sns.kdeplot(final_train["Age"][final_train.Survived == 0], color="lightcoral", shade=True)
4 plt.legend(['Survived', 'Died'])
5 plt.title('Density Plot of Age for Surviving Population and Deceased Population')
6 ax.set(xlabel='Age')
7 plt.xlim(-10,85)
8 plt.show()

```




In [81]:

```

1 plt.figure(figsize=(20,8))
2 avg_survival_byage = final_train[["Age", "Survived"]].groupby(['Age'], as_index=False).mean()
3 g = sns.barplot(x='Age', y='Survived', data=avg_survival_byage, color="LightSeaGreen")
4 plt.show()

```




In [82]:

```
1 final_train['IsMinor']=np.where(final_train['Age']<=16, 1, 0)
2 print(final_train['IsMinor'])
```

```
16    1
17    0
18    0
19    0
20    0
21    0
22    1
23    0
24    1
25    0
26    0
27    0
28    0
29    0
30    0
31    0
32    0
33    0
34    0
35    0
```

In [83]:

```
1 sns.barplot(x='TravelAlone', y='Survived', data=final_train, color="mediumturquoise")
2 plt.show()
```



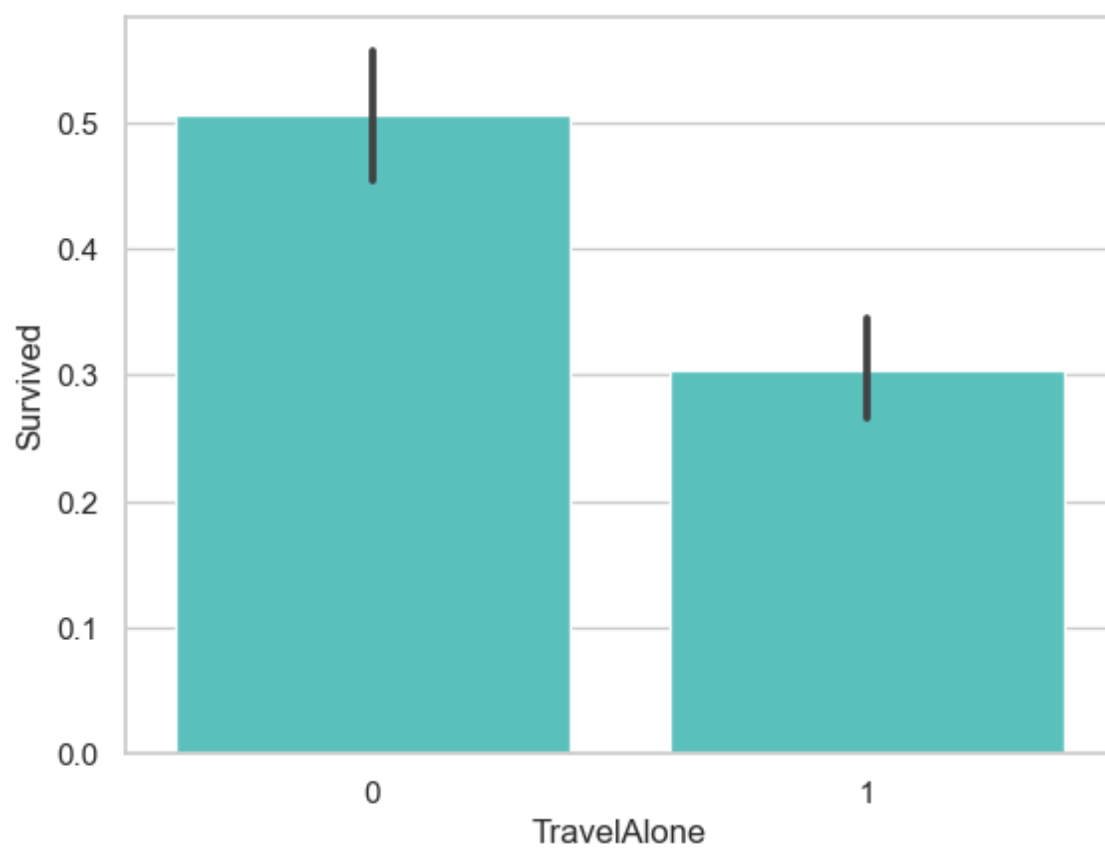
In [84]:

```
1 final_test['IsMinor']=np.where(final_test['Age']<=16, 1, 0)
2 print(final_test['IsMinor'])
```

```
802 1
803 1
804 0
805 0
806 0
807 0
808 0
809 0
810 0
811 0
812 0
813 1
814 0
815 0
816 0
817 0
818 0
819 1
820 0
821 0
```

In [86]:


```
1 sns.barplot(x='TravelAlone', y='Survived', data=final_train, color="mediumturquoise")
2 plt.show()
```



In [87]:



```
1 import seaborn as sns
2 import matplotlib.pyplot as plt
3 # Assuming 'train_df' is your DataFrame containing the data
4 sns.barplot(x='Sex', y='Survived', data=train_df, color='aquamarine')
5 plt.show()
```



In []:



1